

# Hybrid CNN-Transformer Waste Classifier

Report - Assignment 2

Daniele Gotti - 1079011

December 2025

## 1 Introduction

In this project, I developed a multi-class image classifier to categorize waste objects into nine distinct classes (e.g., Plastic, Paper, Metal). The core challenge was to implement a hybrid architecture combining a Convolutional Neural Network (CNN) backbone for feature extraction with a Transformer Encoder for global context processing. Given a training dataset of images, I employed cross-validation strategies to optimize the model's hyperparameters and ensure generalization. The final model was trained on the full dataset and uploaded to HuggingFace for deployment.

## 2 Methodology

### 2.1 Data Preprocessing and Augmentation

The dataset consists of images divided into 9 classes. To enable hyperparameter tuning, I split the provided data into a training set and a validation set (typically utilizing an 80/20 split). All images were resized to  $224 \times 224$  pixels and normalized using ImageNet statistics. To combat overfitting, a significant challenge observed in early experiments, I implemented an aggressive data augmentation pipeline for the training phase, which included:

- Random horizontal flips ( $p=0.5$ );
- Random rotations (up to 15 degrees);
- Color jittering (brightness and contrast).

### 2.2 Hybrid Model Architecture

I implemented a class named `ResnetTransformer` that fuses a CNN with a Transformer:

1. Backbone: I used a ResNet-18, removing the last two layers. This outputs a feature map of shape  $512 \times 7 \times 7$ .
2. Tokenization: the spatial features are flattened into 49 tokens ( $7 \times 7$ ). A linear projection layer maps these from 512 dimensions to a customizable embedding dimension.
3. Transformer: I prepended a learnable CLS token and added positional embeddings. These inputs are processed by a standard PyTorch `TransformerEncoder`.
4. Classification: the final logits are computed by passing the output of the CLS token through a LayerNorm followed by a Linear head.

### 3 Hyperparameter Tuning and Analysis

I conducted a systematic series of experiments (labeled Experiment 0 through 6, followed by 3 Final Tests) to identify the optimal configuration. The accompanying notebook documents the code and plots for each step.

#### 3.1 Baseline and Complexity Analysis

I established a baseline (Exp 0) using a frozen backbone, 128 embedding dimension, and 1 transformer layer.

- Baseline results: the model reached a plateau around 70% accuracy with signs of overfitting.
- Increasing capacity (Exp 1-3): I attempted to increase the model capacity by doubling the layers (to 2 or 4) and the embedding dimension (to 256). Contrary to expectations, increasing complexity worsened overfitting, widening the gap between training and validation performance without improving validation accuracy.
- Attention heads (Exp 4): increasing the number of heads from 4 to 8 had a negligible impact on performance.

#### 3.2 Impact of Regularization and Fine-Tuning

Since architectural changes alone did not solve the overfitting, I shifted focus to data and training dynamics.

- Data augmentation (Exp 5): introducing the augmentation pipeline described in Section 2.1 successfully eliminated overfitting, causing the training and validation curves to overlap. However, the accuracy remained limited by the frozen backbone.
- Unfreezing the backbone (Exp 6): unfreezing the ResNet-18 weights allowed the CNN to learn features specific to the waste dataset. This was the most critical improvement, significantly boosting accuracy.

#### 3.3 Refining the Training Process

The final phase focused on balancing stability and learning speed.

- Learning Rate: I tested lower learning rates ( $1e - 5$  vs  $5e - 5$ ). While  $1e - 5$  produced extremely stable curves, it led to underfitting within the epoch limit. A value of  $3e - 5$  provided the best balance.
- Epochs: I extended training to 20 epochs to allow the fine-tuned model to converge.

## 4 Selection of the Optimal Model

Based on the analysis, I selected the configuration from "Final Test 3" as the optimal model. Although it showed a slight gap between training and validation curves compared to the most conservative tests, it achieved the highest absolute validation accuracy (~87%), maximizing predictive power.

The winning configuration is:

- Backbone: ResNet-18 (unfrozen).
- Transformer: 1 layer, 8 heads, 256 embedding dimension.

- Data augmentation: active (flip, rotation, jitter).
- Optimizer: Adam with learning rate  $3e - 5$  (and exponential scheduler).
- Training epochs: 20.

## 5 Conclusion

I successfully implemented a hybrid CNN-Transformer classifier for waste detection. My analysis demonstrated that for this specific dataset, simply increasing the transformer's size was ineffective. Instead, the combination of aggressive data augmentation and fine-tuning the convolutional backbone proved to be the decisive factors in reaching high performance.

For the final deployment, I retrained the optimal model on the complete dataset to maximize its knowledge base. The final model weights were saved and uploaded to the HuggingFace repository `daniele-gotti/Waste_Classifier`.

## 6 Contributions

I completed this project individually and was responsible for all aspects of the work, including code implementation, experimental analysis, and report writing.

## 7 AI Usage Acknowledgment

I acknowledge that no AI software was used for generating the implementation code or for writing the core content of this report. I used an AI-based tool solely for the purpose of refining English grammar and ensuring a scientific writing style.